


# Kernel methods

CSE 250B

# Deviations from linear separability

## Noise




Find a separator that minimizes  
a convex loss function related  
to the number of mistakes.


e.g. SVM, logistic regression.

# Deviations from linear separability

## Noise



## Systematic deviation




Find a separator that minimizes  
a convex loss function related  
to the number of mistakes.

e.g. SVM, logistic regression.


What to do with this?

## Adding new features



Actual boundary is something like  $x_1 = x_2^2 + 5$ .

## Adding new features




Actual boundary is something like  $x_1 = x_2^2 + 5$ .

- This is quadratic in  $x = (x_1, x_2)$
- But it is linear in  $\Phi(x) = (x_1, x_2, x_1^2, x_2^2, x_1x_2)$

**Basis expansion:** embed data in higher-dimensional feature space.  
Then we can use a linear classifier!

# Basis expansion for quadratic boundaries

How to deal with a  
**quadratic** boundary?



Idea: augment the regular features  $x = (x_1, x_2, \dots, x_d)$  with

$$x_1^2, x_2^2, \dots, x_d^2$$

$$x_1x_2, x_1x_3, \dots, x_{d-1}x_d$$

Enhanced data vectors of the form:

$$\Phi(x) = (x_1, \dots, x_d, x_1^2, \dots, x_d^2, x_1x_2, \dots, x_{d-1}x_d)$$

## Quick question

Suppose  $x = (x_1, x_2, x_3)$ . What is the dimension of  $\Phi(x)$ ?

Suppose  $x = (x_1, \dots, x_d)$ . What is the dimension of  $\Phi(x)$ ?




# Perceptron revisited


Learning in the higher-dimensional feature space:

- $w = 0$  and  $b = 0$
- while some  $y(w \cdot \Phi(x) + b) \leq 0$  :
  - $w = w + y \Phi(x)$
  - $b = b + y$


# Perceptron with basis expansion: examples




## Perceptron with basis expansion: examples



# Perceptron with basis expansion: examples



## Perceptron with basis expansion: examples



# Perceptron with basis expansion

Learning in the higher-dimensional feature space:

- $w = 0$  and  $b = 0$
- while some  $y(w \cdot \Phi(x) + b) \leq 0$  :
  - $w = w + y \Phi(x)$
  - $b = b + y$

# Perceptron with basis expansion

Learning in the higher-dimensional feature space:

- $w = 0$  and  $b = 0$
- while some  $y(w \cdot \Phi(x) + b) \leq 0$  :
  - $w = w + y \Phi(x)$
  - $b = b + y$

**Problem:** number of features has now increased dramatically.  
For MNIST, with quadratic boundary: from 784 to 308504.

# Perceptron with basis expansion

Learning in the higher-dimensional feature space:

- $w = 0$  and  $b = 0$
- while some  $y(w \cdot \Phi(x) + b) \leq 0$  :
  - $w = w + y \Phi(x)$
  - $b = b + y$

**Problem:** number of features has now increased dramatically.  
For MNIST, with quadratic boundary: from 784 to 308504.

**The kernel trick: implement this without ever writing down a vector in the higher-dimensional space!**



# The kernel trick

- $w = 0$  and  $b = 0$
- while some  $y^{(i)}(w \cdot \Phi(x^{(i)}) + b) \leq 0$  :
  - $w = w + y^{(i)} \Phi(x^{(i)})$
  - $b = b + y^{(i)}$

# The kernel trick

- $w = 0$  and  $b = 0$
- while some  $y^{(i)}(w \cdot \Phi(x^{(i)}) + b) \leq 0$  :
  - $w = w + y^{(i)} \Phi(x^{(i)})$
  - $b = b + y^{(i)}$

① Represent  $w$  in **dual** form:  $\alpha = (\alpha_1, \dots, \alpha_n)$ .

$$w = \sum_{j=1}^n \alpha_j y^{(j)} \Phi(x^{(j)})$$

# The kernel trick

- $w = 0$  and  $b = 0$
- while some  $y^{(i)}(w \cdot \Phi(x^{(i)}) + b) \leq 0$  :
  - $w = w + y^{(i)} \Phi(x^{(i)})$
  - $b = b + y^{(i)}$

① Represent  $w$  in **dual** form:  $\alpha = (\alpha_1, \dots, \alpha_n)$ .

$$w = \sum_{j=1}^n \alpha_j y^{(j)} \Phi(x^{(j)})$$

② Compute  $w \cdot \Phi(x)$  using the dual representation.

$$w \cdot \Phi(x) = \sum_{j=1}^n \alpha_j y^{(j)} (\Phi(x^{(j)}) \cdot \Phi(x))$$

# The kernel trick

- $w = 0$  and  $b = 0$
- while some  $y^{(i)}(w \cdot \Phi(x^{(i)}) + b) \leq 0$  :
  - $w = w + y^{(i)} \Phi(x^{(i)})$
  - $b = b + y^{(i)}$

- ① Represent  $w$  in **dual** form:  $\alpha = (\alpha_1, \dots, \alpha_n)$ .

$$w = \sum_{j=1}^n \alpha_j y^{(j)} \Phi(x^{(j)})$$

- ② Compute  $w \cdot \Phi(x)$  using the dual representation.

$$w \cdot \Phi(x) = \sum_{j=1}^n \alpha_j y^{(j)} (\Phi(x^{(j)}) \cdot \Phi(x))$$

- ③ Compute  $\Phi(x) \cdot \Phi(z)$  without ever writing out  $\Phi(x)$  or  $\Phi(z)$ .

# Computing dot products

First, in 2-d.

Suppose  $x = (x_1, x_2)$  and  $\Phi(x) = (x_1, x_2, x_1^2, x_2^2, x_1x_2)$ .

Actually, tweak a little:  $\Phi(x) = (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, x_2^2, \sqrt{2}x_1x_2)$

What is  $\Phi(x) \cdot \Phi(z)$ ?

# Computing dot products

Suppose  $x = (x_1, x_2, \dots, x_d)$  and

$$\Phi(x) = (1, \sqrt{2}x_1, \dots, \sqrt{2}x_d, x_1^2, \dots, x_d^2, \sqrt{2}x_1x_2, \dots, \sqrt{2}x_{d-1}x_d)$$

$$\begin{aligned}\Phi(x) \cdot \Phi(z) &= (1, \sqrt{2}x_1, \dots, \sqrt{2}x_d, x_1^2, \dots, x_d^2, \sqrt{2}x_1x_2, \dots, \sqrt{2}x_{d-1}x_d) \cdot \\ &\quad (1, \sqrt{2}z_1, \dots, \sqrt{2}z_d, z_1^2, \dots, z_d^2, \sqrt{2}z_1z_2, \dots, \sqrt{2}z_{d-1}z_d) \\ &= 1 + 2 \sum_i x_i z_i + \sum_i x_i^2 z_i^2 + 2 \sum_{i \neq j} x_i x_j z_i z_j \\ &= (1 + x_1 z_1 + \dots + x_d z_d)^2 = (1 + x \cdot z)^2\end{aligned}$$

For MNIST:

We are computing dot products in 308504-dimensional space.  
But it takes time proportional to 784, the original dimension!

# Kernel Perceptron

Learning from data  $(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)}) \in \mathcal{X} \times \{-1, 1\}$

**Primal form:**

- $w = 0$  and  $b = 0$
- while there is some  $i$  with  $y^{(i)}(w \cdot \Phi(x^{(i)}) + b) \leq 0$  :
  - $w = w + y^{(i)} \Phi(x^{(i)})$
  - $b = b + y^{(i)}$

**Dual form:**  $w = \sum_j \alpha_j y^{(j)} \Phi(x^{(j)})$ , where  $\alpha \in \mathbb{R}^n$

- $\alpha = 0$  and  $b = 0$
- while some  $i$  has  $y^{(i)} \left( \sum_j \alpha_j y^{(j)} \Phi(x^{(j)}) \cdot \Phi(x^{(i)}) + b \right) \leq 0$  :
  - $\alpha_i = \alpha_i + 1$
  - $b = b + y^{(i)}$

To classify a new point  $x$ :  $\text{sign} \left( \sum_j \alpha_j y^{(j)} \Phi(x^{(j)}) \cdot \Phi(x) + b \right)$ .

## Does this work with SVMs?

$$\begin{aligned} \text{(PRIMAL)} \quad & \min_{w \in \mathbb{R}^d, b \in \mathbb{R}, \xi \in \mathbb{R}^n} \quad \|w\|^2 + C \sum_{i=1}^n \xi_i \\ \text{s.t.:} \quad & y^{(i)}(w \cdot x^{(i)} + b) \geq 1 - \xi_i \quad \text{for all } i = 1, 2, \dots, n \\ & \xi \geq 0 \end{aligned}$$

$$\begin{aligned} \text{(DUAL)} \quad & \max_{\alpha \in \mathbb{R}^n} \quad \sum_{i=1}^n \alpha_i - \sum_{i,j=1}^n \alpha_i \alpha_j y^{(i)} y^{(j)} (x^{(i)} \cdot x^{(j)}) \\ \text{s.t.:} \quad & \sum_{i=1}^n \alpha_i y^{(i)} = 0 \\ & 0 \leq \alpha_i \leq C \end{aligned}$$

Solution:  $w = \sum_i \alpha_i y^{(i)} x^{(i)}$ .



# Kernel SVM

- 1 **Basis expansion.** Mapping  $x \mapsto \Phi(x)$ .
- 2 **Learning.** Solve the dual problem:

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^n} \quad & \sum_{i=1}^n \alpha_i - \sum_{i,j=1}^n \alpha_i \alpha_j y^{(i)} y^{(j)} (\Phi(x^{(i)}) \cdot \Phi(x^{(j)})) \\ \text{s.t.:} \quad & \sum_{i=1}^n \alpha_i y^{(i)} = 0 \\ & 0 \leq \alpha_i \leq C \end{aligned}$$


This yields  $w = \sum_i \alpha_i y^{(i)} \Phi(x^{(i)})$ . Offset  $b$  also follows.

- 3 **Classification.** Given a new point  $x$ , classify as


$$\text{sign} \left( \sum_i \alpha_i y^{(i)} (\Phi(x^{(i)}) \cdot \Phi(x)) + b \right).$$

# Kernel Perceptron vs. Kernel SVM: examples

Perceptron:




SVM:




# Kernel Perceptron vs. Kernel SVM: examples

Perceptron:




SVM:




# Polynomial decision boundaries

When decision surface is a polynomial of order  $p$ :



# Polynomial decision boundaries


When decision surface is a polynomial of order  $p$ :



- Let  $\Phi(x)$  consist of all terms of order  $\leq p$ , such as  $x_1 x_2^2 x_3^{p-3}$ .  
(How many such terms are there, roughly?)

# Polynomial decision boundaries


When decision surface is a polynomial of order  $p$ :



- Let  $\Phi(x)$  consist of all terms of order  $\leq p$ , such as  $x_1 x_2^2 x_3^{p-3}$ .  
(How many such terms are there, roughly?)
- Degree- $p$  polynomial in  $x \Leftrightarrow$  linear in  $\Phi(x)$ .

# Polynomial decision boundaries


When decision surface is a polynomial of order  $p$ :



- Let  $\Phi(x)$  consist of all terms of order  $\leq p$ , such as  $x_1 x_2^2 x_3^{p-3}$ .  
(How many such terms are there, roughly?)
- Degree- $p$  polynomial in  $x \Leftrightarrow$  linear in  $\Phi(x)$ .
- Same trick works:  $\Phi(x) \cdot \Phi(z) = (1 + x \cdot z)^p$ .

# Polynomial decision boundaries

When decision surface is a polynomial of order  $p$ :



- Let  $\Phi(x)$  consist of all terms of order  $\leq p$ , such as  $x_1 x_2^2 x_3^{p-3}$ .  
(How many such terms are there, roughly?)
- Degree- $p$  polynomial in  $x \Leftrightarrow$  linear in  $\Phi(x)$ .
- Same trick works:  $\Phi(x) \cdot \Phi(z) = (1 + x \cdot z)^p$ .
- **Kernel function:**  $k(x, z) = (1 + x \cdot z)^p$ .



# String kernels

Sequence data:

- text documents
- speech signals
- protein sequences

# String kernels

## Sequence data:

- text documents
- speech signals
- protein sequences

Each data point is a sequence of arbitrary length. This yields input spaces like:

$$\mathcal{X} = \{A, C, G, T\}^*$$

$$\mathcal{X} = \{\text{English words}\}^*$$

# String kernels

## Sequence data:

- text documents
- speech signals
- protein sequences

Each data point is a sequence of arbitrary length. This yields input spaces like:

$$\mathcal{X} = \{A, C, G, T\}^*$$

$$\mathcal{X} = \{\text{English words}\}^*$$

What kind of embedding  $\Phi(x)$  is suitable for variable-length sequences  $x$ ?

# String kernels

Sequence data:

- text documents
- speech signals
- protein sequences

Each data point is a sequence of arbitrary length. This yields input spaces like:

$$\mathcal{X} = \{A, C, G, T\}^*$$

$$\mathcal{X} = \{\text{English words}\}^*$$

What kind of embedding  $\Phi(x)$  is suitable for variable-length sequences  $x$ ?

We will use an infinite-dimensional embedding!

## String kernels, cont'd

For each substring  $s$ , define *feature*:

$$\Phi_s(x) = \# \text{ of times substring } s \text{ appears in } x$$

and let  $\Phi(x)$  be a vector with one coordinate for each string:

$$\Phi(x) = (\Phi_s(x) : \text{all strings } s).$$

## String kernels, cont'd

For each substring  $s$ , define *feature*:

$$\Phi_s(x) = \# \text{ of times substring } s \text{ appears in } x$$

and let  $\Phi(x)$  be a vector with one coordinate for each string:

$$\Phi(x) = (\Phi_s(x) : \text{all strings } s).$$

Example: the embedding of “aardvark” includes features

$$\Phi_{\text{ar}}(\text{aardvark}) = 2, \quad \Phi_{\text{th}}(\text{aardvark}) = 0, \dots$$

Linear classifier based on such features is very expressive.

## String kernels, cont'd

For each substring  $s$ , define *feature*:

$$\Phi_s(x) = \# \text{ of times substring } s \text{ appears in } x$$

and let  $\Phi(x)$  be a vector with one coordinate for each string:

$$\Phi(x) = (\Phi_s(x) : \text{all strings } s).$$

Example: the embedding of “aardvark” includes features

$$\Phi_{\text{ar}}(\text{aardvark}) = 2, \quad \Phi_{\text{th}}(\text{aardvark}) = 0, \dots$$

Linear classifier based on such features is very expressive.

To compute  $k(x, z) = \Phi(x) \cdot \Phi(z)$ :

*for each substring  $s$  of  $x$ : count how often  $s$  appears in  $z$*

## String kernels, cont'd

For each substring  $s$ , define *feature*:

$$\Phi_s(x) = \# \text{ of times substring } s \text{ appears in } x$$

and let  $\Phi(x)$  be a vector with one coordinate for each string:

$$\Phi(x) = (\Phi_s(x) : \text{all strings } s).$$

Example: the embedding of “aardvark” includes features

$$\Phi_{ar}(\text{aardvark}) = 2, \quad \Phi_{th}(\text{aardvark}) = 0, \dots$$

Linear classifier based on such features is very expressive.

To compute  $k(x, z) = \Phi(x) \cdot \Phi(z)$ :

*for each substring  $s$  of  $x$ : count how often  $s$  appears in  $z$*

Using dynamic programming, this takes time  $O(|x| \cdot |z|)$ .



# The kernel function

We never explicitly construct the embedding  $\Phi(x)$ .

- What we actually use: **kernel function**  $k(x, z) = \Phi(x) \cdot \Phi(z)$ .
- Think of  $k(x, z)$  as a **measure of similarity** between  $x$  and  $z$ .
- Rewrite learning algorithm and final classifier in terms of  $k$ .

# The kernel function

We never explicitly construct the embedding  $\Phi(x)$ .

- What we actually use: **kernel function**  $k(x, z) = \Phi(x) \cdot \Phi(z)$ .
- Think of  $k(x, z)$  as a **measure of similarity** between  $x$  and  $z$ .
- Rewrite learning algorithm and final classifier in terms of  $k$ .

## Kernel Perceptron:

- $\alpha = 0$  and  $b = 0$
- while some  $i$  has  $y^{(i)} \left( \sum_j \alpha_j y^{(j)} k(x^{(j)}, x^{(i)}) + b \right) \leq 0$ :
  - $\alpha_i = \alpha_i + 1$
  - $b = b + y^{(i)}$

To classify a new point  $x$ :  $\text{sign} \left( \sum_j \alpha_j y^{(j)} k(x^{(j)}, x) + b \right)$ .

# Kernel SVM, revisited

- 1 **Kernel function.** Define a similarity function  $k(x, z)$ .
- 2 **Learning.** Solve the dual problem:

$$\begin{aligned} \max_{\alpha \in \mathbb{R}^n} \quad & \sum_{i=1}^n \alpha_i - \sum_{i,j=1}^n \alpha_i \alpha_j y^{(i)} y^{(j)} k(x^{(i)}, x^{(j)}) \\ \text{s.t.:} \quad & \sum_{i=1}^n \alpha_i y^{(i)} = 0 \\ & 0 \leq \alpha_i \leq C \end{aligned}$$

This yields  $\alpha$ . Offset  $b$  also follows.

- 3 **Classification.** Given a new point  $x$ , classify as

$$\text{sign} \left( \sum_i \alpha_i y^{(i)} k(x^{(i)}, x) + b \right).$$

# Choosing the kernel function

The final classifier is a **similarity-weighted vote**,

$$F(x) = \alpha_1 y^{(1)} k(x^{(1)}, x) + \cdots + \alpha_n y^{(n)} k(x^{(n)}, x)$$

(plus an offset term,  $b$ ).

# Choosing the kernel function

The final classifier is a **similarity-weighted vote**,

$$F(x) = \alpha_1 y^{(1)} k(x^{(1)}, x) + \cdots + \alpha_n y^{(n)} k(x^{(n)}, x)$$

(plus an offset term,  $b$ ).

Can we choose  $k$  to be **any** similarity function?

# Choosing the kernel function

The final classifier is a **similarity-weighted vote**,

$$F(x) = \alpha_1 y^{(1)} k(x^{(1)}, x) + \cdots + \alpha_n y^{(n)} k(x^{(n)}, x)$$

(plus an offset term,  $b$ ).

Can we choose  $k$  to be **any** similarity function?

- Not quite: need  $k(x, z) = \Phi(x) \cdot \Phi(z)$  for *some* embedding  $\Phi$ .

# Choosing the kernel function

The final classifier is a **similarity-weighted vote**,

$$F(x) = \alpha_1 y^{(1)} k(x^{(1)}, x) + \cdots + \alpha_n y^{(n)} k(x^{(n)}, x)$$

(plus an offset term,  $b$ ).

Can we choose  $k$  to be **any** similarity function?

- Not quite: need  $k(x, z) = \Phi(x) \cdot \Phi(z)$  for *some* embedding  $\Phi$ .
- **Mercer's condition**: same as requiring that for any finite set of points  $x^{(1)}, \dots, x^{(m)}$ , the  $m \times m$  similarity matrix  $K$  given by

$$K_{ij} = k(x^{(i)}, x^{(j)})$$

is positive semidefinite.

# The RBF kernel


A popular similarity function: the **Gaussian kernel** or **RBF kernel**

$$k(x, z) = e^{-\|x-z\|^2/s^2},$$


where  $s$  is an adjustable scale parameter.




## RBF kernel: examples




## RBF kernel: examples



## RBF kernel: examples



## RBF kernel: examples



# The scale parameter

Recall prediction function:

$$F(x) = \alpha_1 y^{(1)} k(x^{(1)}, x) + \cdots + \alpha_n y^{(n)} k(x^{(n)}, x).$$

For the RBF kernel,  $k(x, z) = e^{-\|x-z\|^2/s^2}$ ,

- 1 How does this function behave as  $s \uparrow \infty$ ?
- 2 How does this function behave as  $s \downarrow 0$ ?
- 3 As we get more data, should we increase or decrease  $s$ ?

# Kernels: postscript

## ① Customized kernels

- For different domains (NLP, biology, speech, ...)
- Over different structures (sequences, sets, graphs, ...)

# Kernels: postscript

## ① Customized kernels

- For different domains (NLP, biology, speech, ...)
- Over different structures (sequences, sets, graphs, ...)

## ② Learning the kernel function

Given a set of plausible kernels, find a linear combination of them that works well.

# Kernels: postscript

## ① Customized kernels

- For different domains (NLP, biology, speech, ...)
- Over different structures (sequences, sets, graphs, ...)

## ② Learning the kernel function

Given a set of plausible kernels, find a linear combination of them that works well.

## ③ Speeding up learning and prediction

The  $n \times n$  kernel matrix  $k(x^{(i)}, x^{(j)})$  is a bottleneck for large  $n$ .

One idea:

- Go back to the primal space!
- Replace the embedding  $\Phi$  by a low-dimensional mapping  $\tilde{\Phi}$  such that

$$\tilde{\Phi}(x) \cdot \tilde{\Phi}(z) \approx \Phi(x) \cdot \Phi(z).$$

This can be done, for instance, by writing  $\Phi$  in the Fourier basis and then randomly sampling features.