

SUV Color Space & Filtering

Computer Vision I
CSE252A
Lecture 9

CS252A, Fall 2017

Computer Vision I

Announcement

- HW2 assigned

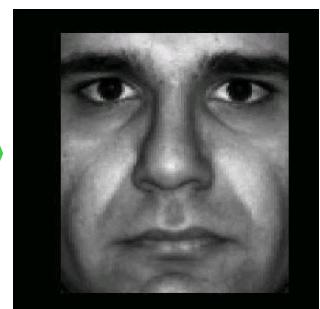
SUV Color Space for Photometric Stereo of Glossy Objects

Note: You'll use this in HW2.

CS252A, Fall 2017

Computer Vision I

Lambertian Photometric Stereo



Reconstruction with Albedo Map

Computer Vision I

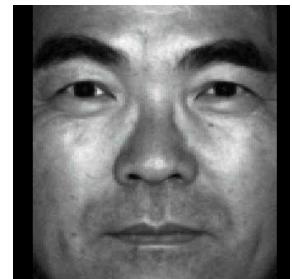
Without the albedo map



CS252A, Fall 2017

Computer Vision I

Another person



CS252A, Fall 2017

Computer Vision I

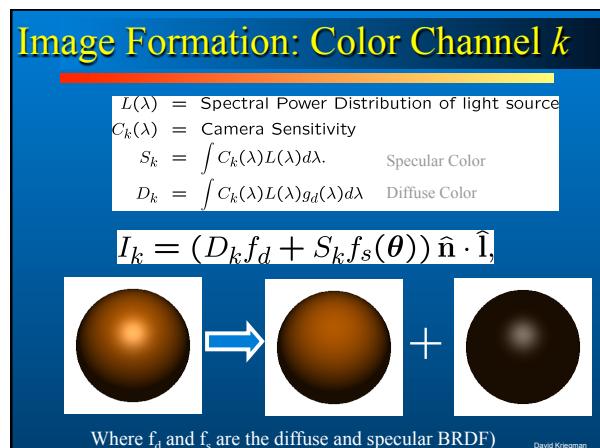
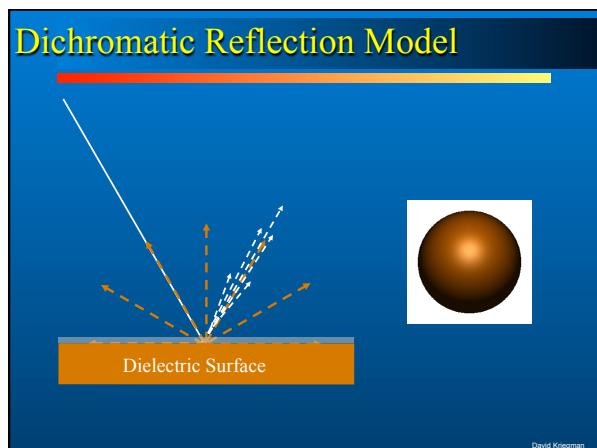
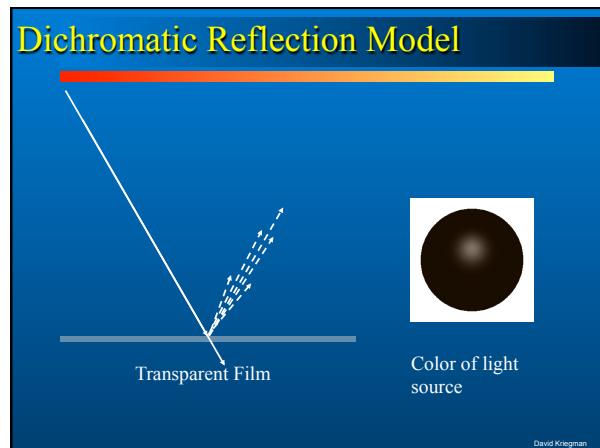
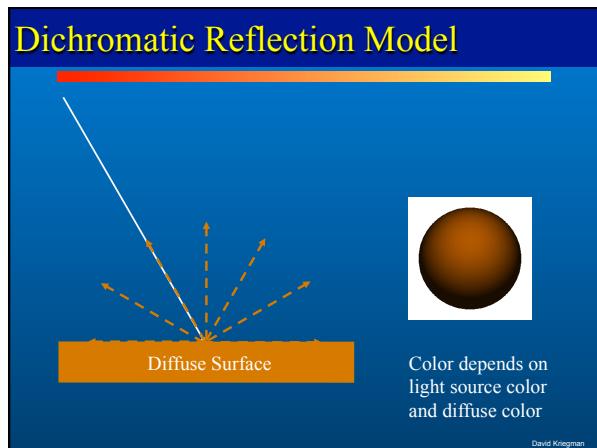
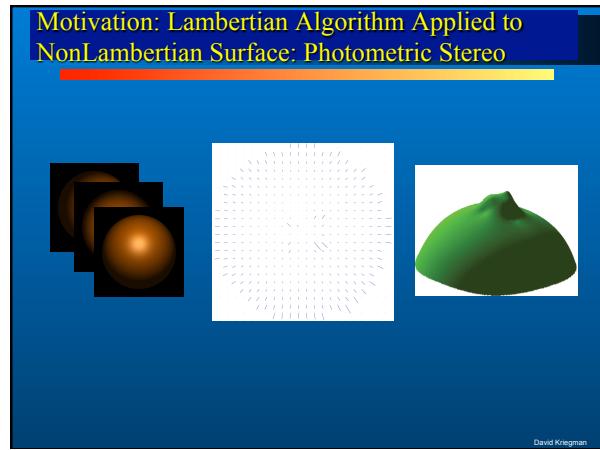
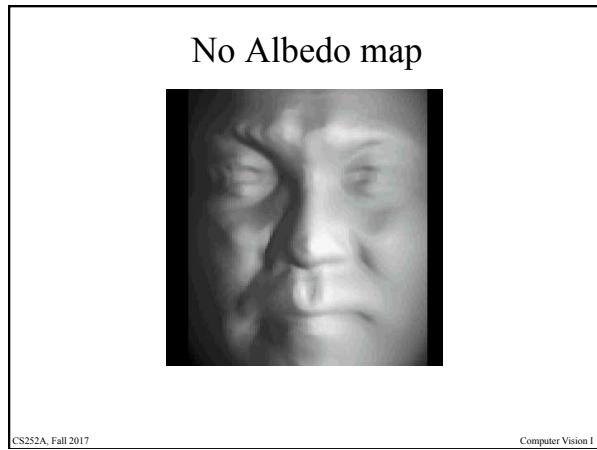
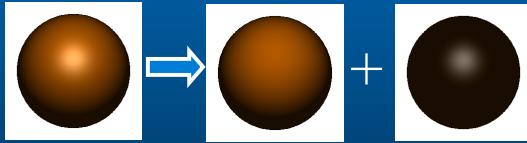


Image formation: Color Channel k

$$I_k = (D_k f_d + S_k f_s(\theta)) \hat{n} \cdot \hat{l},$$

$$\begin{bmatrix} I_r \\ I_g \\ I_b \end{bmatrix} = (f_d \mathbf{n} \cdot \hat{l}) \begin{bmatrix} D_r \\ D_g \\ D_b \end{bmatrix} + (f_s(\theta) \mathbf{n} \cdot \hat{l}) \begin{bmatrix} S_r \\ S_g \\ S_b \end{bmatrix}$$

Image color lies in span of diffuse color \mathbf{D} and specular color \mathbf{S} .



David Kriegman

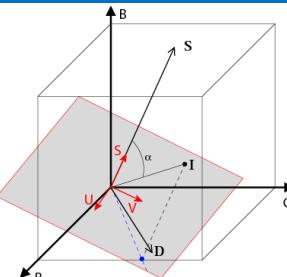
Data-dependent SUV Color Space

$$\mathbf{I}_{SUV} = [R]\mathbf{I}_{RGB}$$

$$[R] = [\mathbf{S} | \mathbf{U} | \mathbf{V}]^T$$

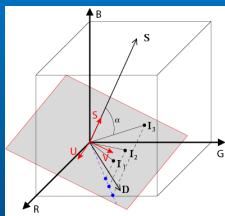
$$[R] \in SO(3)$$

First row of R is specular color \mathbf{S} . Other rows are orthogonal to \mathbf{S}



David Kriegman

Properties of SUV



- Data-dependent.
- Rotational (hence, linear) Transformation.
- The S channel encodes the entire specular component and an unknown amount of diffuse component.
- Shading information is preserved in u and v channels.

$$I_U = \mathbf{r}_2^\top \mathbf{D} f_d \hat{n} \cdot \hat{l}$$

$$I_V = \mathbf{r}_3^\top \mathbf{D} f_d \hat{n} \cdot \hat{l}$$

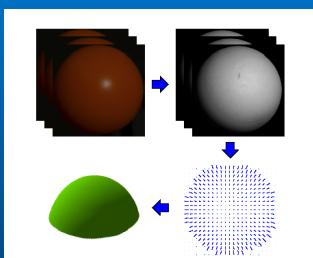
David Kriegman

Example



David Kriegman

Multi-channel Photometric Stereo



David Kriegman

Multi-channel Photometric Stereo

$$\mathbf{J} = [I_U \ I_V]^\top$$

\mathbf{J}^k : 2-channel color vector under the k^{th} light source.

$\hat{\mathbf{l}}^k$: The k^{th} three light source directions.

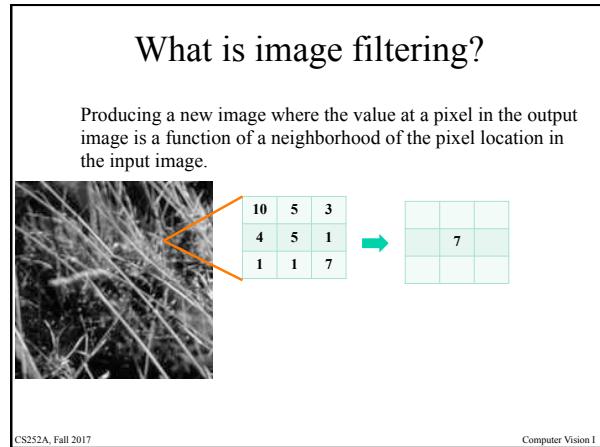
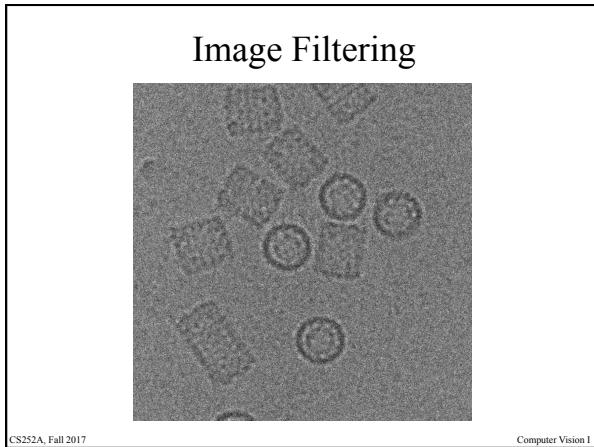
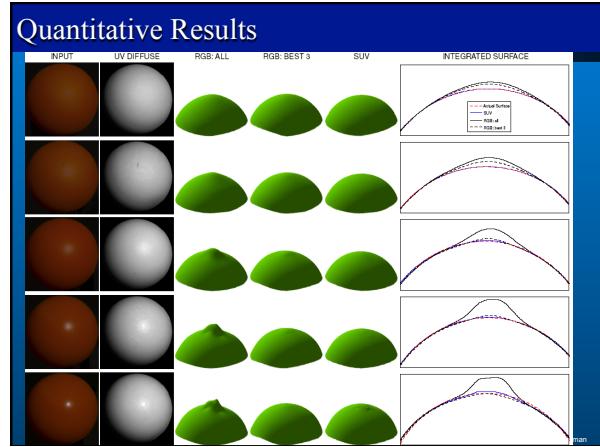
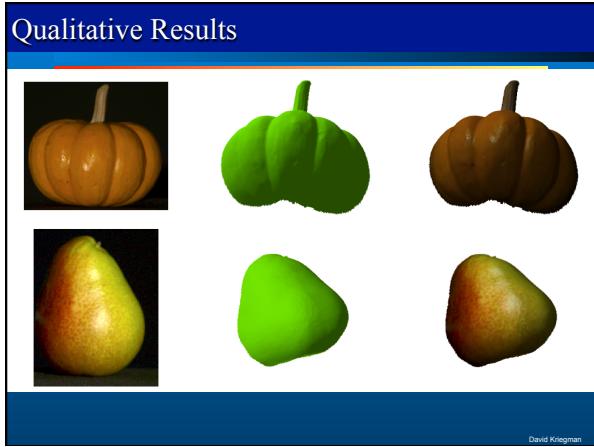
ρ : 2-channel UV albedo.

$$\mathbf{J}^k = [I_U^k \ I_V^k]^\top = (\hat{n} \cdot \hat{\mathbf{l}}^k) \rho,$$

$$\begin{aligned} \text{Shading vector } \mathbf{F} &= [f^1, f^2, f^3]^\top = [\hat{l}^1 \ \hat{l}^2 \ \hat{l}^3]^\top \hat{n} \\ \text{Intensity matrix } \mathbf{J} &= \begin{bmatrix} J_1^1 & J_2^1 \\ J_1^2 & J_2^2 \\ J_1^3 & J_2^3 \end{bmatrix} = \begin{bmatrix} f^1 \rho_U & f^1 \rho_V \\ f^2 \rho_U & f^2 \rho_V \\ f^3 \rho_U & f^3 \rho_V \end{bmatrix} = \mathbf{F} \rho^\top. \end{aligned}$$

The least squares estimate of the shading vector \mathbf{F} is the principal eigenvector of $[\mathbf{J}][\mathbf{J}]^\top$. Once the shading vector is known, the surface normal is found by solving the matrix equation $\mathbf{F} = [\hat{l}^1 \ \hat{l}^2 \ \hat{l}^3]^\top \hat{n}$.

David Kriegman



- ### Linear Filters
- General process:
 - Form new image whose pixels are a weighted sum of original pixel values, using the same set of weights at each point.
 - Properties
 - Output is a linear function of the input
 - Output is a shift-invariant function of the input (i.e. shift the input image two pixels to the left, the output is shifted two pixels to the left)
 - Example: smoothing by averaging
 - form the average of pixels in a neighborhood
 - Example: smoothing with a Gaussian
 - form a weighted average of pixels in a neighborhood
 - Example: finding a derivative
- CS252A, Fall 2017 Computer Vision I

Linear functions

- Simplest: linear filtering.
 - Replace each pixel by a linear combination of its neighbors.
- The prescription for the linear combination is called the “convolution kernel”.

Local image data	$\begin{bmatrix} 10 & 5 & 3 \\ 4 & 1 & 1 \\ 1 & 1 & 7 \end{bmatrix}$	kernel weights	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$	Modified image data	$\begin{bmatrix} 7 \end{bmatrix}$
------------------	--	----------------	---	---------------------	-----------------------------------

(Freeman)

CS252A, Fall 2017

Computer Vision I

Convolution

* $\begin{bmatrix} 1 & 2 & 1 \\ -1 & -2 & -1 \end{bmatrix}$

Kernel (K)

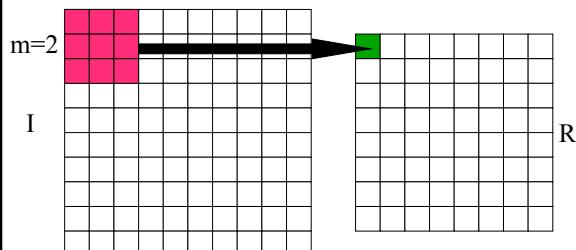
Note: Typically Kernel is relatively small in vision applications.

Image (I)

CS252A, Fall 2017

Computer Vision I

Convolution: $R = K * I$



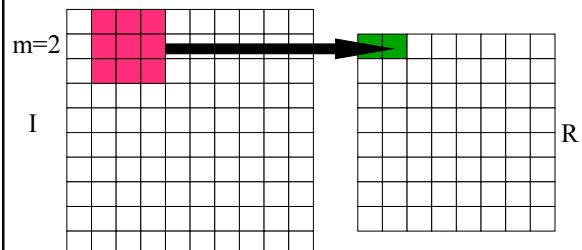
Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$

CS252A, Fall 2017

Computer Vision I

Convolution: $R = K * I$



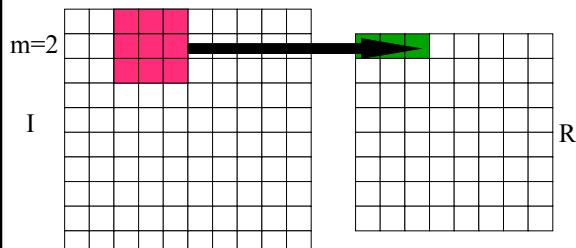
Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$

CS252A, Fall 2017

Computer Vision I

Convolution: $R = K * I$



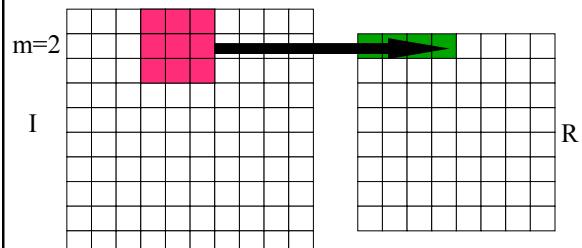
Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$

CS252A, Fall 2017

Computer Vision I

Convolution: $R = K * I$

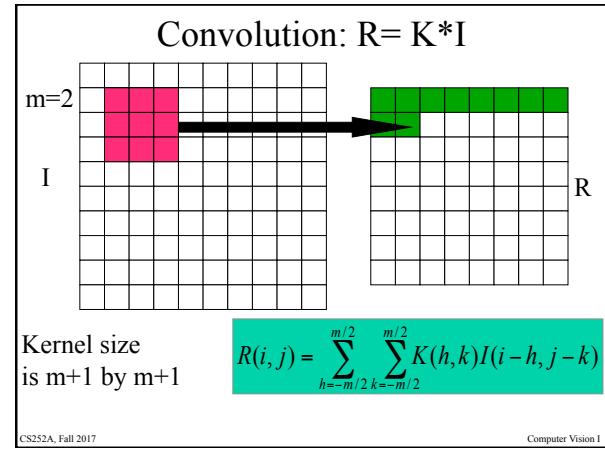
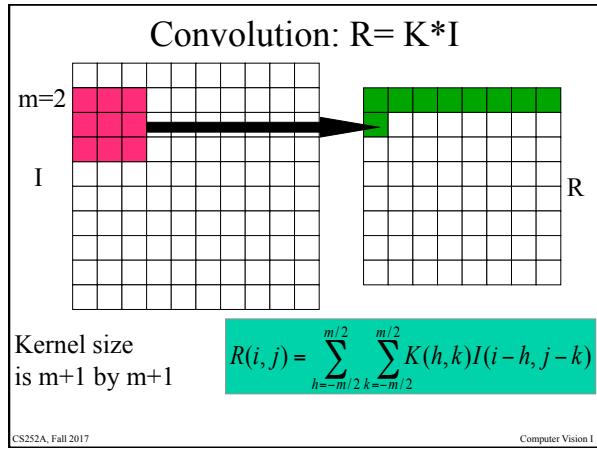
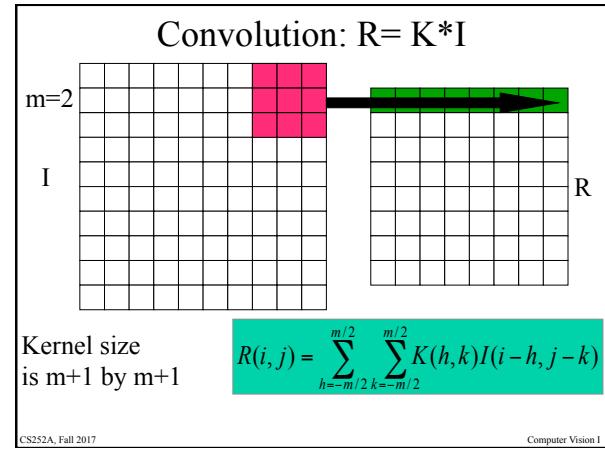
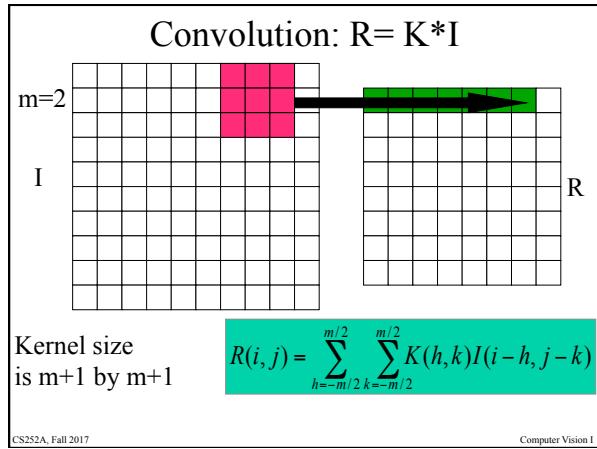
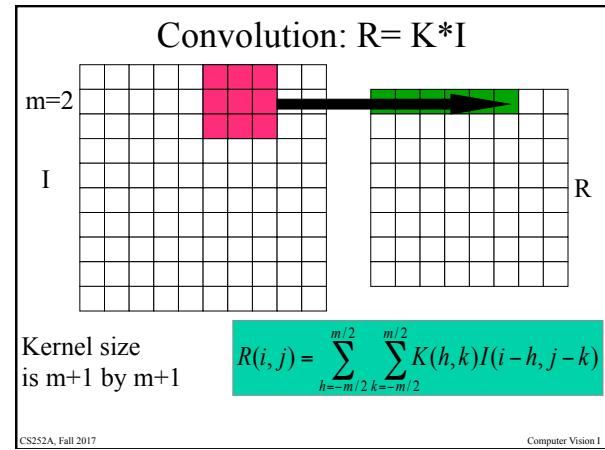
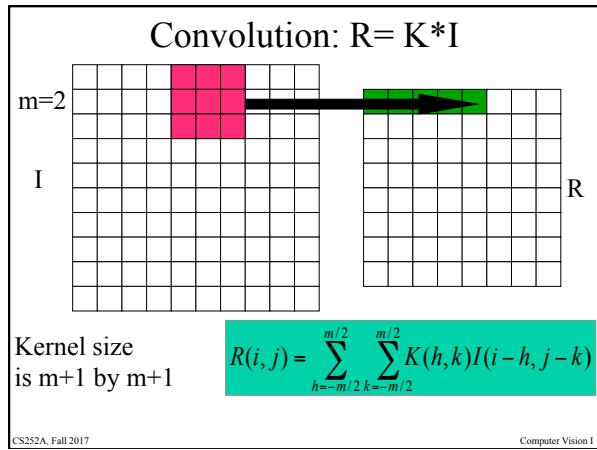


Kernel size
is $m+1$ by $m+1$

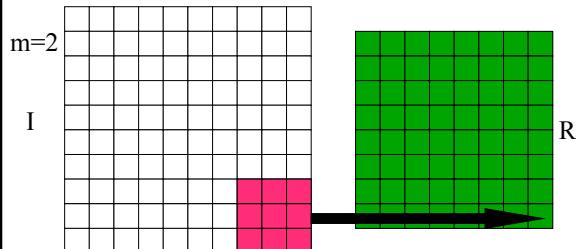
$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$

CS252A, Fall 2017

Computer Vision I



Convolution: $R = K * I$



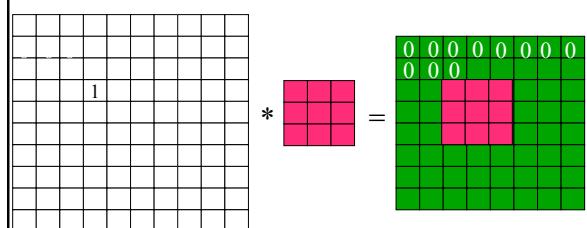
Kernel size
is $m+1$ by $m+1$

$$R(i, j) = \sum_{h=-m/2}^{m/2} \sum_{k=-m/2}^{m/2} K(h, k) I(i-h, j-k)$$

CS252A, Fall 2017

Computer Vision I

Impulse Response



CS252A, Fall 2017

Computer Vision I

Convolutional Neural Networks

- Core operation is, not surprisingly, convolution.
- During training of a CNN, the weights of the convolution kernels are learned.

CS252A, Fall 2017

Computer Vision I

Linear filtering (warm-up slide)



original

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$

?

CS252A, Fall 2017

Computer Vision I

Linear filtering (warm-up slide)



original

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$



Filtered
(no change)

CS252A, Fall 2017

(Swiped from Bill Freeman)

Computer Vision I

Linear filtering



original

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{matrix}$$

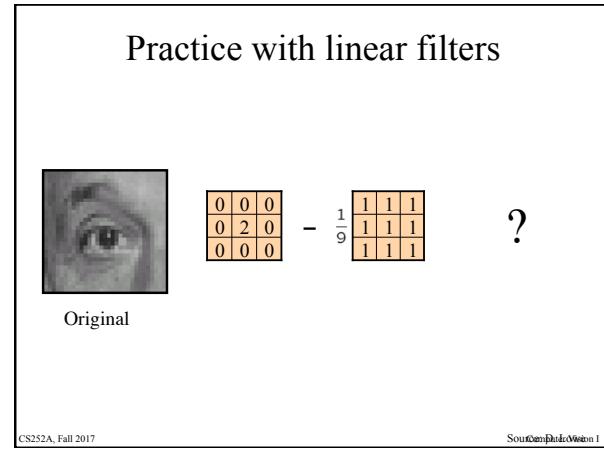
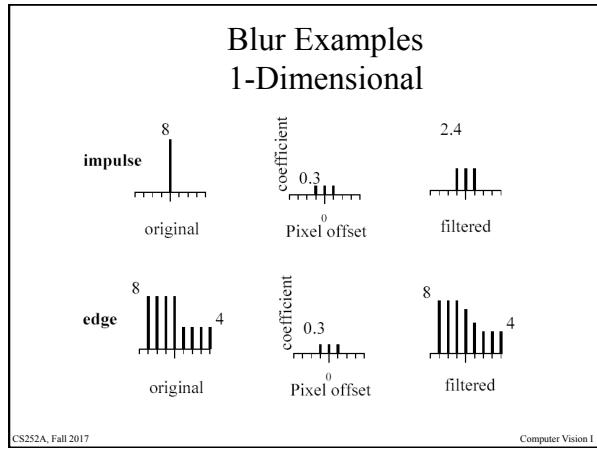
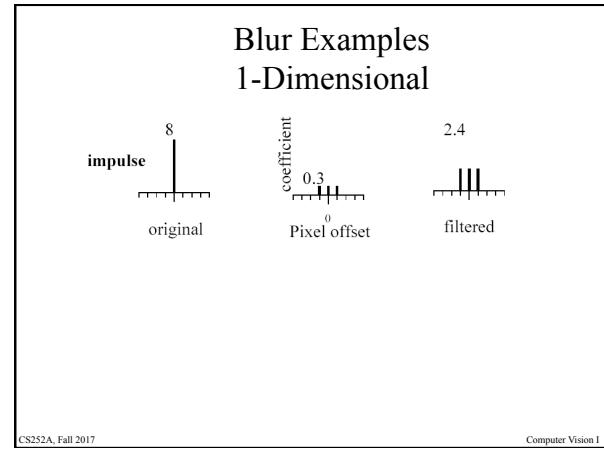
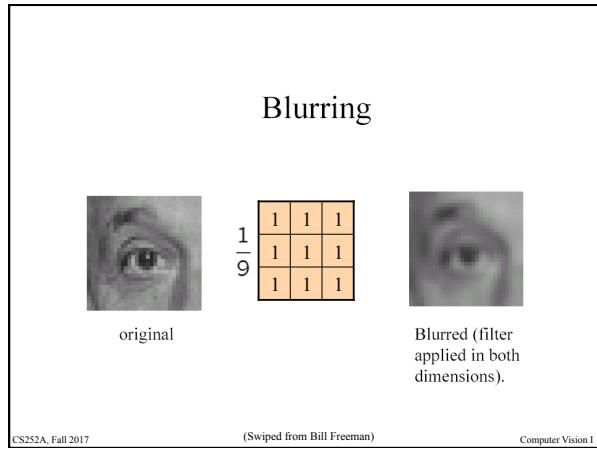
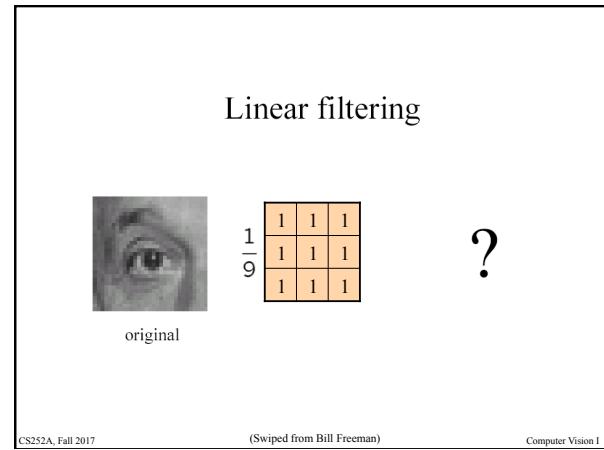
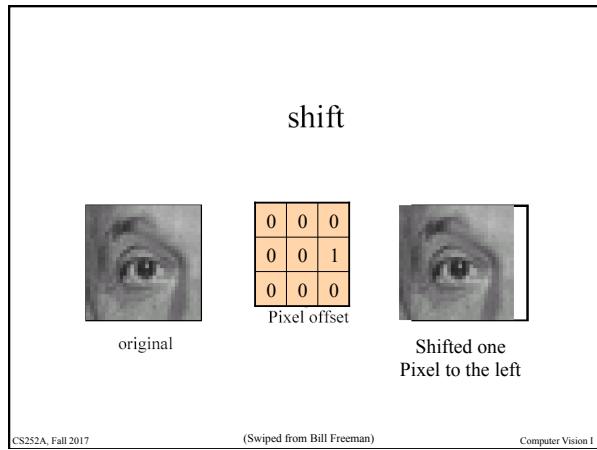
Pixel offset

?

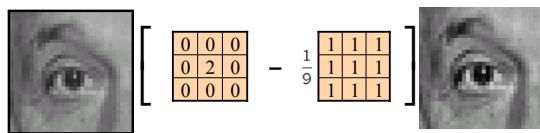
CS252A, Fall 2017

(Swiped from Bill Freeman)

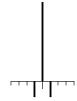
Computer Vision I



Practice with linear filters



Original

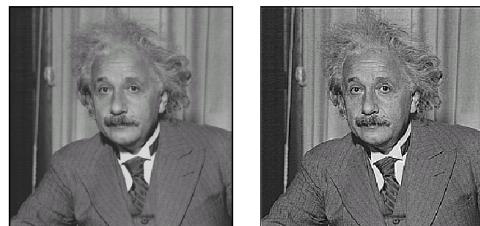


Sharpening filter
- Accentuates differences with local average

Source: Computer Vision I

CS252A, Fall 2017

Sharpening



before

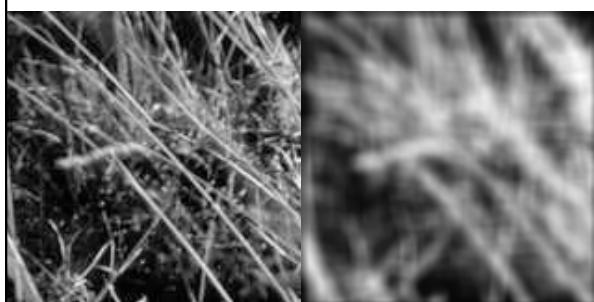
after

CS252A, Fall 2017

Computer Vision I

Smoothing by Averaging

Kernel: 

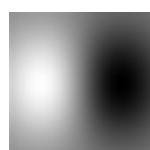


CS252A, Fall 2017

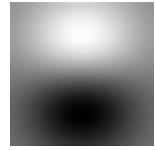
Computer Vision I

Filters are templates

- Applying a filter at some image location can be seen as taking a dot-product between a neighborhood and the kernel.
- Filtering the image is like taking a dot product at each location.
- Insight
 - filters look like the effects they are intended to find
 - filters find effects they look like



CS252A, Fall 2017



Computer Vision I

Properties of Continuous Convolution (Holds for discrete too)

Let f, g, h be images and $*$ denote convolution

$$f * g(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x-u, y-v) g(u, v) du dv$$

- Commutative: $f * g = g * f$
- Associative: $f * (g * h) = (f * g) * h$
- Linear: for scalars a & b and images f, g, h
 $(af + bg) * h = a(f * h) + b(g * h)$
- Differentiation rule

$$\frac{\partial}{\partial x}(f * g) = \frac{\partial f}{\partial x} * g = f * \frac{\partial g}{\partial x}$$

CS252A, Fall 2017

Computer Vision I

Filtering to reduce noise

- Noise is what we're not interested in.
 - We'll discuss simple, low-level noise today:
Light fluctuations; Sensor noise; Quantization effects; Finite precision
 - Not complex: shadows; extraneous objects.
- A pixel's neighborhood contains information about its intensity.
- Averaging noise reduces its effect.

CS252A, Fall 2017

Computer Vision I

Additive noise

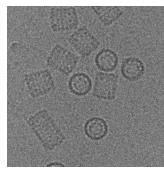
- $I = S + N$. Noise doesn't depend on signal.
- We'll consider:

$$I_i = s_i + n_i \text{ with } E(n_i) = 0$$

s_i deterministic. n_i a random var.

n_i, n_j independent for $i \neq j$

n_i, n_j identically distributed



CS252A, Fall 2017

Computer Vision I

Gaussian
Noise:
sigma=1



Computer Vision I

Guassian
Noise:
sigma=16



CS252A, Fall 2017

Computer Vision I

Averaging Filter

- Mask with positive entries, that sum 1.
- Replaces each pixel with an average of its neighborhood.
- If all weights are equal, it is called a BOX filter.

$$\begin{matrix} & & F \\ & 1 & 1 & 1 \\ 1/9 & 1 & 1 & 1 \\ & 1 & 1 & 1 \end{matrix}$$

(Camps)

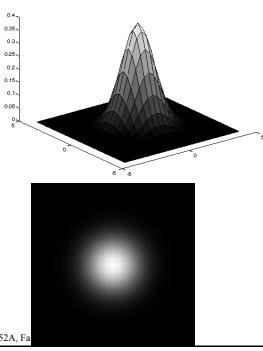
Computer Vision I

An Isotropic Gaussian

- The picture shows a smoothing kernel proportional to

$$e^{-\frac{x^2+y^2}{2\sigma^2}}$$

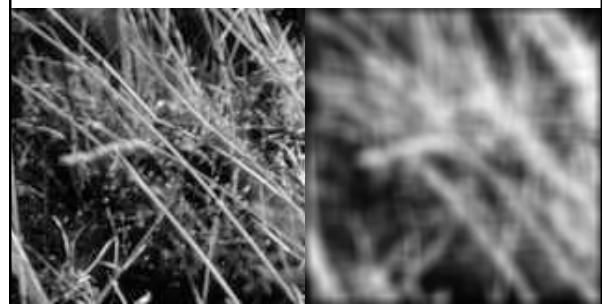
(which is a reasonable model of a circularly symmetric fuzzy blob)



Computer Vision I

Smoothing by Averaging

Kernel:

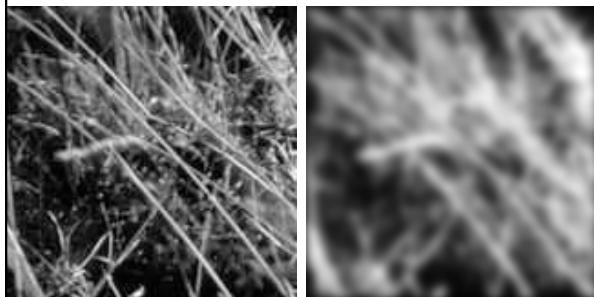


CS252A, Fall 2017

Computer Vision I

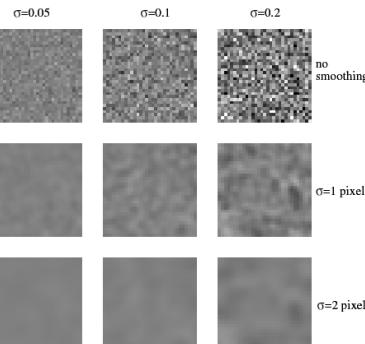
Smoothing with a Gaussian

Kernel: 



CS252A, Fall 2017

Computer Vision I



The effects of smoothing
Each row shows smoothing with gaussians of different width; each column shows different realizations of an image of gaussian noise.

CS252A, Fall 2017

Computer Vision I

Efficient Implementation

- Both, the BOX filter and the Gaussian filter are separable:
 - First convolve each row with a 1D filter
 - Then convolve each column with a 1D filter.

$$\begin{aligned} r(u)*c(v)*I(x,y) &= (r(u)*c(v))*I(u,v) \\ &= (c(u)*r(v))*I(u,v) \\ &= c(u)*r(v)*I(x,y) \end{aligned}$$

CS252A, Fall 2017

Computer Vision I

Fourier Transform

- 1-D transform (signal processing)
- 2-D transform (image processing)
- Consider 1-D

Time domain \leftrightarrow Frequency Domain
Real \leftrightarrow Complex

- Consider time domain signal to be expressed as weighted sum of sinusoids. A sinusoid $\cos(ut+\phi)$ is characterized by its phase ϕ and its frequency u
- The Fourier transform of the signal is a function giving the weights (and phase) as a function of frequency u .

CS252A, Fall 2017

Computer Vision I

Fourier Transform

Discrete Fourier Transform (DFT) of $I[x,y]$

$$F[u, v] \equiv \frac{1}{N} \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} I[x, y] e^{-\frac{2\pi}{N} i (xu + yv)}$$

Inverse DFT

$$I[x, y] \equiv \frac{1}{N} \sum_{u=0}^{N-1} \sum_{v=0}^{N-1} F[u, v] e^{\frac{2\pi}{N} i (ux + vy)}$$

x,y: spatial domain

u,v: frequency domain

Implemented via the "Fast Fourier Transform" algorithm (FFT)

CS252A, Fall 2017

Computer Vision I

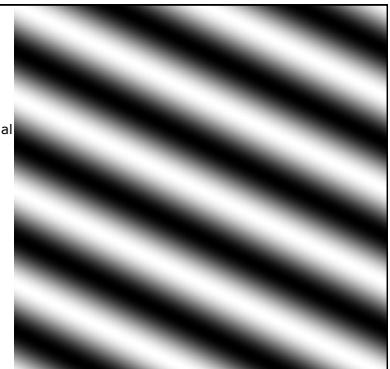
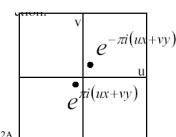
Fourier basis element

$$e^{-i2\pi(ux+vy)}$$

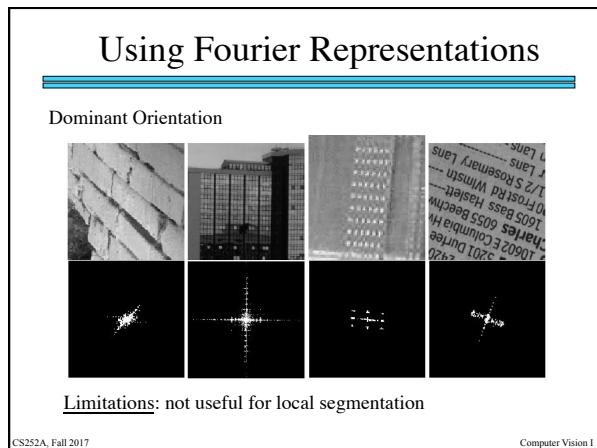
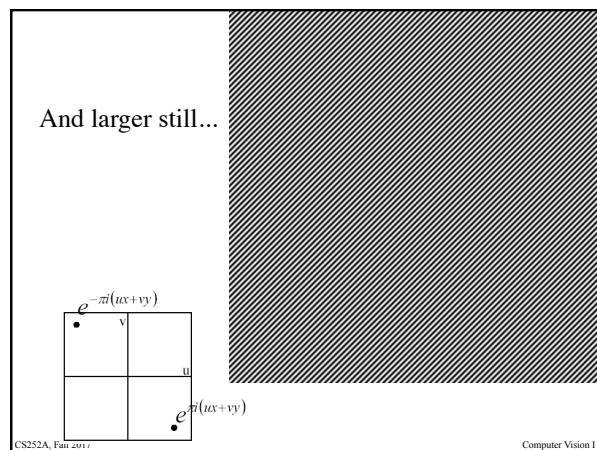
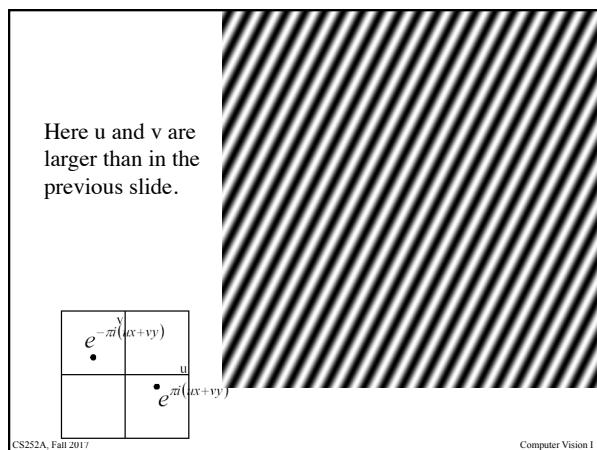
Transform is sum of orthogonal basis functions

Vector (u, v)

- Magnitude gives frequency
- Direction gives orientation.



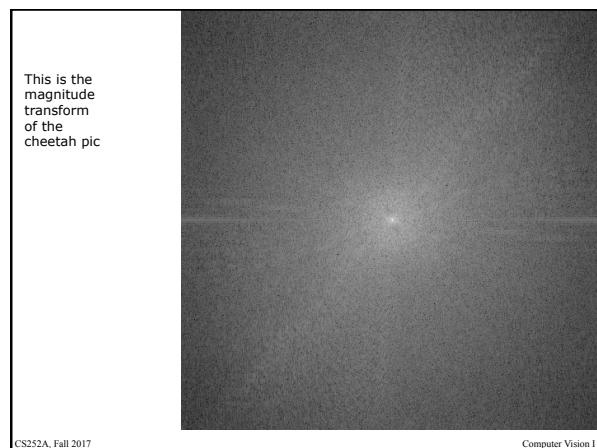
Computer Vision I

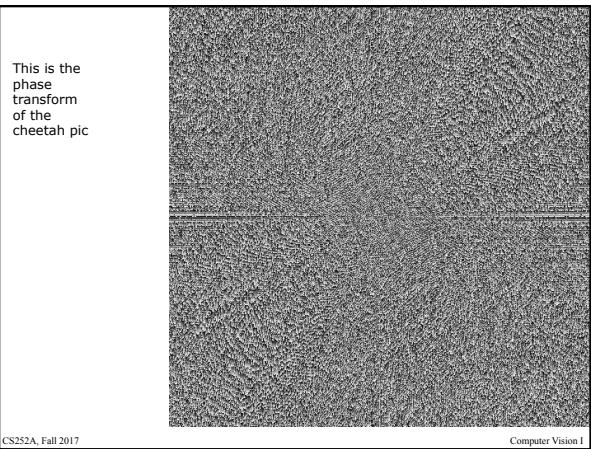


$e^{i\theta} = \cos\theta + i \sin\theta$

- Fourier transform of a real function is complex
 - difficult to plot, visualize
 - instead, we can think of the phase and magnitude of the transform
- Phase is the phase of the complex transform
- Magnitude is the magnitude of the complex transform
- Curious fact
 - all natural images have about the same magnitude transform
 - hence, phase seems to matter, but magnitude largely doesn't
- Demonstration
 - Take two pictures, swap the phase transforms, compute the inverse - what does the result look like?

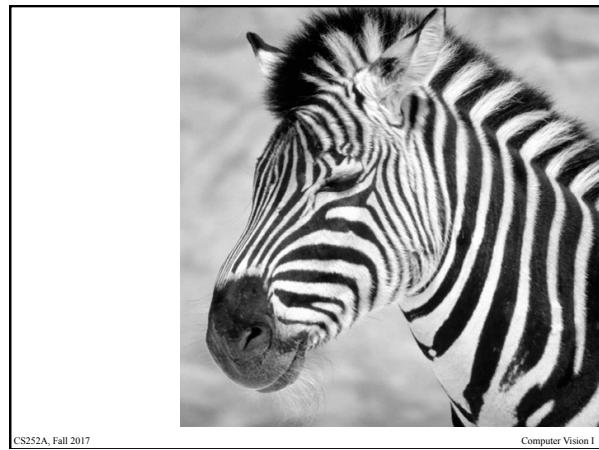
CS252A, Fall 2017 Computer Vision I





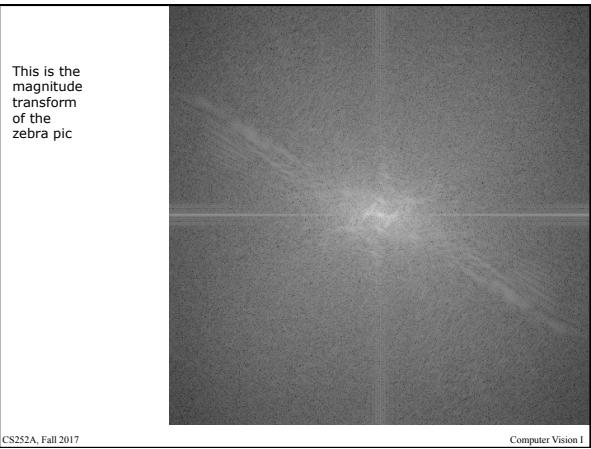
CS252A, Fall 2017

Computer Vision I



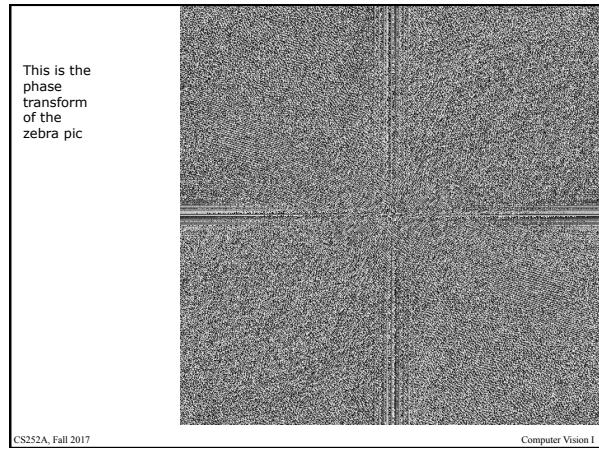
CS252A, Fall 2017

Computer Vision I



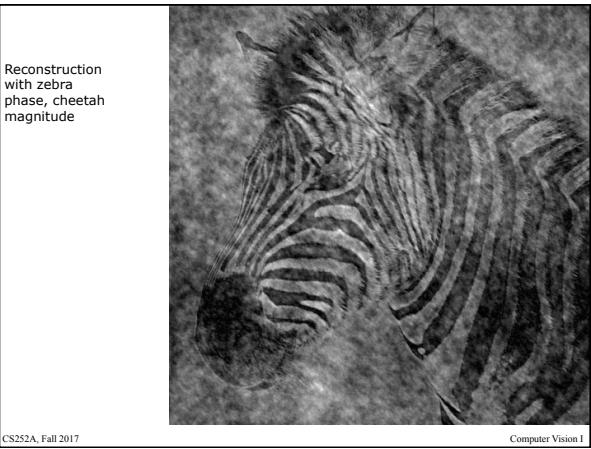
CS252A, Fall 2017

Computer Vision I



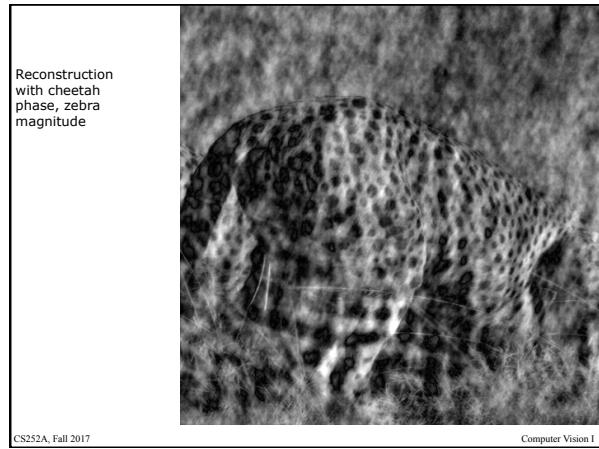
CS252A, Fall 2017

Computer Vision I



CS252A, Fall 2017

Computer Vision I



CS252A, Fall 2017

Computer Vision I

The Fourier Transform and Convolution

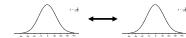
- If H and G are images, and $F(\cdot)$ represents Fourier transform, then
$$F(H^*G) = F(H)F(G)$$
- Thus, one way of thinking about the properties of a convolution is by thinking of how it modifies the frequencies of the image to which it is applied.
- In particular, if we look at the power spectrum, then we see that convolving image H by G attenuates frequencies where G has low power, and amplifies those which have high power.
- This is referred to as the **Convolution Theorem**

CS252A, Fall 2017

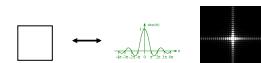
Computer Vision I

Various Fourier Transform Pairs

- Important facts
 - scale function down \Leftrightarrow scale transform up
i.e. high frequency = small details
 - The FT of a Gaussian is a Gaussian.



compare to box function transform



Computer Vision I

Other Types of Noise

- Impulsive noise
 - randomly pick a pixel and randomly set it to a value
 - saturated version is called salt and pepper noise
- Quantization effects
 - Often called noise although it is not statistical
- Unanticipated image structures
 - Also often called noise although it is a real repeatable signal.

CS252A, Fall 2017

Computer Vision I

Some other useful filtering techniques

- Median filter
- Anisotropic diffusion

CS252A, Fall 2017

Computer Vision I

Median filters : principle

Method :

1. rank-order neighbourhood intensities
 2. take middle value
- non-linear filter
 - no new grey levels emerge...

CS252A, Fall 2017

Computer Vision I

Median filters: Example for window size of 3

1,1,1,7,1,1,1



? , 1,1,1,1,1,?

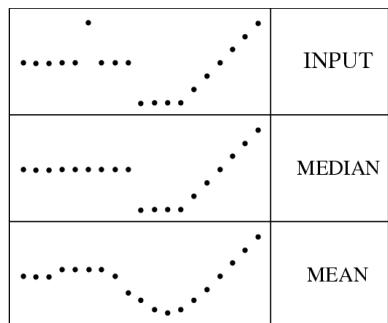
advantage of this type of filter is that it Eliminates spikes (salt & pepper noise).

CS252A, Fall 2017

Computer Vision I

Median filters : example

filters have width 5 :



CS252A, Fall 2017

Computer Vision I

Median filters : analysis

median completely discards the spike,
linear filter always responds to all aspects

median filter preserves discontinuities,
linear filter produces rounding-off effects

DON'T become all too optimistic

Computer Vision I

Median filter : images

3 x 3 median filter :



sharpens edges, destroys edge cusps
and protrusions

CS252A, Fall 2017

Computer Vision I

Median filters : Gauss revisited

Comparison with Gaussian :



e.g. upper lip smoother, eye better preserved

CS252A, Fall 2017

Computer Vision I

Example of median

10 times 3 X 3 median



patchy effect
important details lost (e.g. ear-ring)

CS252A, Fall 2017

Computer Vision I